


When the treatment changes spread or shape, mean- and rank-based tests have essentially no power

Two-sample tests at the 5% level; 5,000 replicates per cell (Monte-Carlo SE = 0.7 percentage points).

Dashed line marks the nominal 5% level; all rank-based tests use exact distribution-free critical values.



Kolmogorov-Smirnov omitted from the panels for legibility (it is uniformly below Anderson-Darling here) but retained in results/power_summary.csv. "Smooth k" = smooth test with k components fixed in advance; "data-driven" selects k = 4 by Schwarz's rule. In the spread-only panel the t-test and Wilcoxon curves coincide at the nominal level.