

RESEARCH ARTICLE

Quantifying Effect Modifier Similarity for Regional Pooling in Multi-Regional Clinical Trials

Author One¹ | Author Two² | Author Three¹

¹Department Name, Institution Name, State, Country

²Department Name, Institution Name, State, Country

Correspondence

Corresponding Author, Department Name, Institution Name, Address.
Email: corresponding@institution.edu

Abstract

The ICH E17 guideline describes regional pooling in multi-regional clinical trials (MRCTs) in terms of the similarity of effect modifier distributions, but provides no quantitative methodology. We propose a planning-stage framework that quantifies effect modifier dissimilarity using the Wasserstein-1 (W_1) distance, retained in the original measurement units of each effect modifier. Unlike the standardized mean difference (SMD), the W_1 distance captures differences in scale and skewness in addition to location—distributional features that may drive treatment effect heterogeneity through non-linear effect modification. Because it requires only baseline effect modifier distributions, W_1 can be computed from any representative data source, including prior trials, disease registries, and real-world evidence (RWE) databases—enabling a common planning scenario in which a small-sample region identifies suitable pooling partners before the trial begins. We develop a dual-pathway clinical calibration: when prior evidence on conditional average treatment effect (CATE) sensitivity (L) is available, the W_1 distance is translated into a maximum potential regional treatment effect difference ($\Delta_{\max}$) with bootstrap confidence intervals; when L is unknown, the framework reverse-calculates the required CATE sensitivity (L^*) at pre-specified clinical margins. Simulation studies across seven systematic scenarios demonstrated satisfactory bias and coverage of the percentile bootstrap estimator for non-negligible distributional differences at moderate sample sizes per region. We illustrate the framework through a hypothetical thrombolytic MRCT grounded in publicly available GUSTO-I individual patient data, with one region designated as a small-sample anchor and the remaining regions evaluated as pooling partners on two candidate effect modifiers (age and systolic blood pressure). Partner rankings differed markedly between the two effect modifiers, with several partners ranked similar on one but dissimilar on the other, demonstrating that all candidate effect modifiers must be evaluated jointly. A small subset of partners emerged as leading pooling candidates because they ranked low on both candidate effect modifiers and required clinically implausible CATE sensitivity to produce a meaningful treatment effect difference. The principal contribution of this work is to transform the previously qualitative judgment of “similar enough” into a reproducible quantitative basis for sponsor judgment, filling the methodological gap in ICH E17 implementation and supporting evidence-based regional pooling decisions from any representative population data source.

KEY WORDS

multi-regional clinical trial; ICH E17; effect modifier; Wasserstein distance; regional pooling; distributional similarity

1 | INTRODUCTION

Multi-regional clinical trials (MRCTs), conducted across multiple countries or regulatory regions under a single protocol, have become the standard paradigm for global pharmaceutical development.^{1,2} This approach offers substantial benefits: accelerated

Abbreviations: AMI, acute myocardial infarction; CATE, conditional average treatment effect; CDF, cumulative distribution function; CI, confidence interval; EDF, empirical distribution function; EU, European Union; GUSTO, Global Utilization of Streptokinase and t-PA for Occluded coronary arteries; ICH, International Council for Harmonisation; IPD, individual patient data; IQR, interquartile range; KL, Kullback–Leibler; KS, Kolmogorov–Smirnov; MRCT, multi-regional clinical trial; NMPA, National Medical Products Administration; RWE, real-world evidence; SBP, systolic blood pressure; SMD, standardized mean difference; US, United States.

timelines, broader generalizability, and earlier access to new therapies for patients worldwide. The International Council for Harmonisation (ICH) E17 guideline, adopted in 2017, established principles for planning and designing MRCTs, with a central assumption that treatment effects are generalizable across the target population.³

Several regulatory authorities expect demonstration of treatment effect consistency in regional subpopulations,^{4,5} yet individual regional sample sizes in global MRCTs are often insufficient for such demonstration. A key strategy for addressing this challenge is the *regional pooling approach*, wherein regions with similar patient characteristics are grouped for analysis to enhance the ability to assess regional consistency.³ The ICH E17 guideline describes pooling decisions in terms of the similarity of effect modifier distributions.³

An effect modifier is a baseline patient characteristic—such as age, disease severity, or genetic marker—for which the treatment benefit differs across subgroups. For example, if younger patients respond better to treatment than older patients, age is an effect modifier. When such heterogeneity exists, even if the drug works identically at the individual level, regions with different patient compositions may observe different average treatment effects. A region with predominantly younger patients would show larger benefits than a region with predominantly older patients, not because the drug works differently, but because the patient mix differs. This fundamental relationship underscores why effect modifier distributional similarity is critical to the validity of regional pooling. In this paper, we focus on continuous effect modifiers, such as age and disease severity scores; categorical effect modifiers raise distinct methodological considerations and are outside the present scope.

Despite the regulatory importance of effect modifier distributional similarity, current practice lacks a standardized, clinically calibrated methodology. The ICH E17 guideline provides no specific metric, threshold, or statistical procedure for determining when distributions are “similar enough.” Recent regulatory guidance has highlighted this gap. Song et al., writing from the China NMPA perspective on ICH E17 implementation, note the challenge of operationalizing pooling criteria without quantitative tools.⁵ Long et al. further discuss basic considerations for consistency evaluation under ICH E17.⁶ The most concrete existing proposal is that of Komiya et al., members of the ICH E17 expert working group, who operationalize pooling through cluster analysis on region-level effect modifier summaries, with inter-regional distances measured between representative values and weighted for relevance via the Lasso.⁷ This proposal leaves two gaps that we address. First, as operationalized through a representative value per effect modifier, it characterizes each regional distribution by its location and inherits the location-centred blind spot of the standardized mean difference, rather than capturing the full distribution. Second, it explicitly calls for a judgment of whether effect modifier differences are “clinically meaningful or not” but supplies no metric to operationalize that judgment.

This absence of quantitative tools is particularly acute at the *planning stage*, when sponsors must identify suitable pooling partners for small-sample regions. The pooling rationale requires quantitative evidence that partner regions’ patient populations are “similar enough,” yet no established method exists to provide such evidence. A workshop report by Matsushima et al. documented that regional imbalance in CRP-positive/MRI-negative status contributed to apparent inconsistency in the secukinumab MRCT.⁴ This example illustrates the practical consequences when effect modifier distributional differences are not assessed quantitatively at the planning stage.

Although general-purpose comparison tools such as the standardized mean difference are routinely applied to baseline covariates,⁸ they were not designed for assessing effect modifier distributional similarity across regions. As a single summary measure, the standardized mean difference cannot detect differences in variance or distributional shape. When effect modification is non-linear, these distributional features directly influence regional average treatment effects, and a comparison based on means alone will miss them.

This paper addresses these gaps by proposing a quantitative framework for assessing effect modifier distributional similarity across regions, based on the Wasserstein-1 (W_1) distance between two cumulative distribution functions, retained in the original measurement units of the effect modifier. Geometrically, W_1 equals the area between the two cumulative distribution functions. Unlike the standardized mean difference, W_1 captures differences in variance and shape simultaneously. We estimate W_1 with bootstrap confidence intervals and link it to clinical interpretation through a calibration framework. The Wasserstein distance provides a theoretically grounded measure of distributional dissimilarity in this setting.

The remainder of this paper is organized as follows. Section 2 presents the methodological framework, including the index definition, clinical calibration, and estimation procedure. Section 3 describes a simulation study. Section 4 demonstrates the framework using publicly available individual patient data. Section 5 discusses implications, limitations, and future directions.

2 | METHODS

An effect modifier is a factor for which the treatment effect varies across subgroups. Formally, let the conditional average treatment effect (CATE) be denoted $\tau(x) = E[Y(1) - Y(0) \mid X = x]$, where X is the effect modifier. When this function is non-constant, the average treatment effect observed in region r depends on the distribution F_r of the effect modifier in that region:

$$\bar{\tau}_r = \int \tau(x) dF_r(x). \quad (1)$$

Assessing pooling suitability therefore requires a measure of distributional difference between regions.

2.1 | Existing Approaches and Their Limitations

The most widely used measure for comparing covariate distributions across groups is the standardized mean difference (SMD), defined as

$$\text{SMD} = \frac{\bar{X}_1 - \bar{X}_2}{S_{\text{pooled}}}, \quad (2)$$

where $\bar{X}_r$ denotes the sample mean in region r and S_{pooled} is the pooled standard deviation.⁸ Because treatment effects in clinical trials are typically summarized as mean differences, SMD provides a scale-free comparison directly relevant to the quantities that drive clinical decision-making. However, by construction SMD captures only differences in location. Two distributions with identical means but markedly different variances—for example, $N(50, 5^2)$ and $N(50, 15^2)$ —yield $\text{SMD} = 0$ despite a threefold difference in spread.

Distributional distance measures address this limitation by comparing distributions beyond summary statistics. Empirical distribution function (EDF) tests such as the Kolmogorov–Smirnov (KS) statistic, $D_{\text{KS}} = \sup_x |F_1(x) - F_2(x)|$, and related measures (Cramér–von Mises, Anderson–Darling) quantify differences in the full shape of distributions by comparing CDFs directly. Information-theoretic divergences such as Kullback–Leibler (KL) divergence,

$$D_{\text{KL}}(P \parallel Q) = \int p(x) \log \frac{p(x)}{q(x)} dx, \quad (3)$$

take a density-based approach, measuring how one probability distribution differs from another.

However, each of these approaches has limitations when applied to comparing effect modifier distributions across regions. SMD captures only location differences and is blind to variance and shape. EDF-based statistics such as the KS statistic capture full distributional differences but lack a clinically interpretable scale and have no theoretical connection to treatment effect heterogeneity. KL divergence is asymmetric—comparing region A to region B yields a different value than comparing B to A—can diverge to infinity when empirical distributions do not fully overlap, and requires density estimation that is impractical with small regional samples. Crucially, none of these measures offers a theoretical link that would translate a distributional distance into a bound on regional treatment effect differences. These gaps motivate a measure that (i) captures distributional features beyond location, (ii) provides a clinically interpretable scale, and (iii) offers such a theoretical link. The Wasserstein-1 distance and its clinical calibration, developed in the following subsections, satisfy all three requirements.

2.2 | Per-EM Wasserstein-1 Distance

Among distributional distances that could capture the full distributional differences relevant to regional pooling, we adopt the Wasserstein-1 distance (W_1) for its unique theoretical connection to treatment effect heterogeneity (Proposition 1 below). The W_1 distance (sometimes called the Earth Mover’s Distance) admits three equivalent representations on the real line, each highlighting a different facet of its interpretation.^{9,10}

Let F_1 and F_2 denote the cumulative distribution functions (CDFs) of an effect modifier in two regions. Then

$$W_1(F_1, F_2) = \int_{-\infty}^{\infty} |F_1(x) - F_2(x)| dx = \int_0^1 |F_1^{-1}(u) - F_2^{-1}(u)| du = \inf_{\pi \in \Pi(F_1, F_2)} \int |x - y| d\pi(x, y), \quad (4)$$




FIGURE 1 The Wasserstein-1 distance as the area between cumulative distribution functions, illustrated with a right-skewed Region 1 (gamma distribution) and a symmetric Region 2 (normal distribution). Panel (A): probability densities. Panel (B): cumulative distribution functions; the shaded area equals $W_1(F_1, F_2)$.

where F_r^{-1} is the quantile function and $\Pi(F_1, F_2)$ is the set of probability couplings with marginals F_1 and F_2 . The CDF form expresses W_1 as the area between the two CDFs (Figure 1); the quantile form as the integrated absolute difference of quantile functions; and the Kantorovich form as the minimum expected absolute displacement to transport one distribution onto the other. The three forms are mathematically equivalent.¹⁰ Unlike the standardized mean difference, W_1 responds to differences in variance, skewness, and other distributional features.¹¹

A property of W_1 central to our framework is that it carries the units of the underlying effect modifier. In each of the three equivalent forms in equation (4), the integrand has units of the variable X : the CDF values are dimensionless probabilities, the quantile function F^{-1} returns values in the original measurement units, and the coupling integral sums distances $|x - y|$ in those units. Consequently, $W_1(F_1, F_2)$ is expressed on the original scale of the effect modifier. This dimensional alignment supports direct clinical interpretation: a regional $W_1 = 5$ mmHg in systolic blood pressure means that an expected 5 mmHg of probability mass must be transported between the two distributions to align them. The same logic precludes a meaningful direct comparison of W_1 values across effect modifiers measured on different scales, and we therefore restrict attention in this paper to per-effect-modifier assessment (Section 5 elaborates on this scope choice).

W_1 is a proper metric on the space of probability distributions with finite first moment, satisfying non-negativity, symmetry, identity of indiscernibles, and the triangle inequality.⁹ It is also affine-equivariant in the variable units, satisfying $W_1(F_{aX+b}, F_{aY+b}) = |a| \cdot W_1(F_X, F_Y)$ for any $a \neq 0$. Beyond these basic properties, W_1 admits a Kantorovich–Rubinstein dual representation⁹

$$W_1(F_1, F_2) = \sup_{\|f\|_{\text{Lip}} \leq 1} \left| \int f dF_1 - \int f dF_2 \right|, \quad (5)$$

which expresses W_1 as the largest possible difference in expected values of a 1-Lipschitz function f between the two distributions (the notation $\|f\|_{\text{Lip}} \leq 1$ means $|f(x) - f(y)| \leq |x - y|$ for all x, y). If the conditional average treatment effect function $\tau(x)$ is Lipschitz continuous with constant L_{clinical} , so that $|\tau(x) - \tau(y)| \leq L_{\text{clinical}}|x - y|$ for all x, y , then τ/L_{clinical} is 1-Lipschitz and applying equation (5) yields the central result of this section:

Proposition 1. (Connection to heterogeneity). *If the CATE function is Lipschitz continuous with constant L_{clinical} , then for any two regions with effect modifier distributions F_1 and F_2 ,*

$$|\bar{\tau}_1 - \bar{\tau}_2| \leq L_{\text{clinical}} \cdot W_1(F_1, F_2). \quad (6)$$

Proof. Apply equation (5) with $f = \tau/L_{\text{clinical}}$, which is 1-Lipschitz by assumption. This yields $|\int (\tau/L_{\text{clinical}}) d(F_1 - F_2)| \leq W_1(F_1, F_2)$. Combining with equation (1) for $\bar{\tau}_r$ and rearranging yields equation (6). $\square$

The bound in equation (6) is specific to W_1 . The W_2 distance, while admitting closed-form expressions for Gaussian families, does not possess a Lipschitz dual representation and therefore cannot provide an analogous bound. The alternative measures

discussed in Section 2.1 similarly lack any analogous property.¹² Its right-hand side involves W_1 in its original measurement units, multiplied by a clinical slope L_{clinical} ; no further normalization of W_1 is required to obtain a clinically interpretable bound. We therefore adopt W_1 in its original units as the primary similarity metric throughout this paper. Some prior literature has considered normalizing W_1 by a within-effect-modifier scale measure such as the pooled interquartile range; we show in Section 2.4 and Supplement A that such normalization is mathematically redundant once equation (6) is used for clinical calibration. The Lipschitz assumption underlying Proposition 1 is appropriate when the CATE function varies smoothly with the effect modifier; for threshold or step-function CATEs, L_{clinical} can be arbitrarily large and the bound becomes conservative (Section 5).

2.3 | Estimation

Given samples $\{X_{1,i}\}_{i=1}^{n_1}$ from region 1 and $\{X_{2,j}\}_{j=1}^{n_2}$ from region 2, the Wasserstein-1 distance is estimated using empirical distribution functions. Since $\hat{F}_1$ and $\hat{F}_2$ are piecewise constant (right-continuous step functions), the integral in equation (4) reduces to a finite sum over the intervals defined by the combined order statistics:

$$\hat{W}_1 = \sum_{k=1}^{n_1+n_2-1} |\hat{F}_1(x_{(k)}) - \hat{F}_2(x_{(k)})| \cdot (x_{(k+1)} - x_{(k)}), \quad (7)$$

where $\hat{F}_r$ denotes the empirical CDF (right-continuous) of region r and $x_{(1)} < \dots < x_{(n_1+n_2)}$ are the combined order statistics. The estimator $\hat{W}_1$ inherits the original measurement units of the effect modifier. The asymptotic distribution of W_1 is non-standard: del Barrio et al.¹³ showed that $\sqrt{n} W_1(\hat{F}_n, F) \xrightarrow{d} \int |B(F(x))| dx$, where B is a standard Brownian bridge. Because the quantiles of this limiting distribution depend on the unknown F , no universal critical values are available and direct construction of asymptotic confidence intervals is impractical. We therefore employ the nonparametric percentile bootstrap with $B = 2,000$ replicates.¹⁴ In the two-sample setting, when $F_1 \neq F_2$ the Hadamard derivative of the L_1 functional is linear, so the ordinary nonparametric bootstrap is consistent.¹⁵ Convergence occurs at rate $\sqrt{n_1 n_2 / (n_1 + n_2)}$ (see Appendix A.2 for details). The estimator (7) is computed in $O((n_1 + n_2) \log(n_1 + n_2))$ time, dominated by sorting the combined order statistics.

2.4 | Interpretation and Clinical Calibration

A central feature of W_1 is its connection to treatment effect heterogeneity through Proposition 1. This connection provides the basis for a clinically grounded interpretation framework rather than reliance on fixed numerical cutoffs.

We refer to the Lipschitz constant of the CATE function $\tau(x)$ as the clinical slope L_{clinical} . By definition, L_{clinical} quantifies the maximum change in the treatment effect per unit change in the effect modifier, expressed on the outcome scale of the trial—for example, percentage points of 30-day mortality per year of age, or percentage points per mmHg of systolic blood pressure. When L_{clinical} is estimable from prior data, individual patient data (IPD) meta-regression on treatment-by-covariate interaction provides a per-unit slope directly,^{16,17} and the same quantity is denoted the interaction slope in the methodological literature.^{18,19} More commonly at the MRCT planning stage, however, published subgroup reports document the direction and significance of interaction rather than a per-unit slope, and a precise numerical L_{clinical} is unavailable. We address both situations below.

From Proposition 1, the maximum potential difference in regional average treatment effects attributable to effect modifier distributional differences is

$$\Delta_{\max} = L_{\text{clinical}} \cdot W_1(F_1, F_2), \quad (8)$$

expressed in the units of the trial outcome scale. The formula is invariant to within-effect-modifier scale normalization of W_1 : as shown in Supplement A, any normalized form W_1/X multiplied by the same scale X cancels in equation (8), so adopting W_1 in its original units yields the same clinical-scale calibration as any rescaled version. When L_{clinical} is independently known, equation (8) provides a direct clinical-scale translation of the observed distributional difference. When L_{clinical} is unknown, the relation can be inverted to ask which L_{clinical} would be required for a given W_1 to produce a clinically meaningful $\Delta_{\max}$.

The estimator $\hat{W}_1$ has a percentile bootstrap CI (Section 2.3), and inferential uncertainty propagates linearly to $\Delta_{\max}$ when L_{clinical} is treated as fixed: $[\Delta_{\max,L}, \Delta_{\max,U}] = L_{\text{clinical}} \cdot [W_{1,L}, W_{1,U}]$. When L_{clinical} itself is uncertain or unavailable, we solve equation (8) for L_{clinical} at a pre-specified clinical margin Δ_{clin} (the minimal clinically important treatment effect difference, the

non-inferiority margin, or the targeted overall effect), yielding the required clinical slope

$$L^* = \frac{\Delta_{\text{clin}}}{W_1(F_1, F_2)}. \quad (9)$$

L^* is the smallest clinical slope under which the observed distributional difference would produce a clinically meaningful regional treatment effect difference. If a plausible upper bound on L_{clinical} from prior evidence (denoted L_{UB}) lies below L^* , the distributional difference is unlikely to translate into clinical heterogeneity, regardless of the unknown true L_{clinical} . The two pathways—forward calibration via equation (8) when L_{clinical} is known and reverse comparison via equation (9) against L_{UB} when L_{clinical} is unknown—accommodate the full range of evidence states encountered at the MRCT planning stage.

Specifying L_{clinical} (or its upper bound L_{UB}) draws on disease-area evidence. For the GUSTO-I case study presented in Section 4, the Fibrinolytic Therapy Trialists’ meta-analysis²⁰ and GUSTO-I subgroup analyses^{21,22} together support a conservative upper bound of $L_{\text{UB,age}} \approx 1 \times 10^{-2}$ per year and $L_{\text{UB,SBP}} \approx 2 \times 10^{-3}$ per mmHg on the 30-day mortality scale. “Conservative” here indicates that the chosen bounds exceed the empirical prognostic age slope of approximately 6×10^{-3} per year reported by Lee et al.,²² so that any plausible CATE sensitivity falls below them. Supplement D discusses the derivation of these bounds, alternative specification strategies such as the deft estimator of Fisher et al.,¹⁷ and the limitations of class-level evidence. The methodology does not require an exact L_{clinical} ; sensitivity analyses across a plausible range of values, or comparison of L^* to L_{UB} , accommodate the typical uncertainty in this parameter.

The pooling assessment proceeds by computing per-effect-modifier W_1 with bootstrap CIs across candidate partner regions, deriving Δ_{max} when L_{clinical} is known or L^* at a pre-specified Δ_{clin} when L_{clinical} is unknown, and ranking partners by these quantities. The framework is deliberately estimation-centered rather than testing-based: W_1 has no universal cutoff that separates “similar” from “dissimilar,” and the appropriate threshold depends on the clinical slope L_{clinical} , which is external to the distributional comparison. By presenting per-pair Δ_{max} (or L^*) with bootstrap CIs alongside the clinical context—the overall treatment effect size, the non-inferiority margin, the minimal clinically important difference—the framework provides quantitative inputs for sponsor and regulatory judgment without forcing a binary accept/reject decision.²³

3 | SIMULATION STUDY

We conducted simulation studies to evaluate the estimation properties of the $\hat{W}_1$ estimator: bias, variability, and coverage probability of bootstrap confidence intervals. These properties are essential for the clinical calibration framework, as reliable estimation of W_1 directly determines the quality of Δ_{max} estimates via equation (8). We report operating characteristics for $\hat{W}_1$ in the original units of the effect modifier, with each scenario specified at $\sigma = 10$; bias and CI width are therefore on the W_1 scale, while coverage is unitless and directly comparable across scenarios. We assessed performance across a range of scenarios relevant to MRCT applications.

3.1 | Simulation Design

3.1.1 | Scenarios

We designed seven systematic scenarios for methodological validation, examining controlled distributional differences across location, scale, and combinations thereof. Table 1 summarizes these scenarios with their true W_1 values, in the original units of the effect modifier.

3.1.2 | Simulation Parameters

For each scenario, we generated samples of size $n = 50, 100$, and 200 per region, reflecting sample sizes typical in MRCT regional subgroups. We performed 10,000 replications per scenario–sample size combination to ensure stable estimates of operating characteristics, with Monte Carlo standard errors below 0.005 for all reported proportions. Bootstrap confidence intervals were computed using $B = 2,000$ resamples. All simulations were conducted in R version 4.5.1.

TABLE 1 Simulation scenarios with clinical motivations.

ID	Description	Distribution 1	Distribution 2	True W_1
S1	Null (identical)	$N(50, 10^2)$	$N(50, 10^2)$	0.00
S2	Location 0.2σ	$N(50, 10^2)$	$N(52, 10^2)$	2.00
S3	Location 0.5σ	$N(50, 10^2)$	$N(55, 10^2)$	5.00
S4	Location 1.0σ	$N(50, 10^2)$	$N(60, 10^2)$	10.00
S5	Scale $1.5\times$	$N(50, 10^2)$	$N(50, 15^2)$	3.99
S6	Skew (Log-normal)	$N(50, 10^2)$	$\text{LogN}(\mu, \sigma)$	12.15
S7	Location + Scale	$N(50, 10^2)$	$N(55, 15^2)$	5.85

The Log-normal in S6 is parameterized to match mean 50 with CV $\approx 53\%$, typical of clinical biomarkers such as ALT. S7 combines location and scale differences, representing the most realistic MRCT pattern. True W_1 values for S1–S4 are exact ($W_1 = |\mu_1 - \mu_2|$ for equal-variance normals); S5 uses the closed form $W_1 = \sqrt{2/\pi} |\sigma_1 - \sigma_2|$; S6 and S7 are computed via Monte Carlo integration ($n_{MC} = 10^6$).

3.1.3 | Evaluation Metrics

Consistent with the estimation-centered framework, we evaluated the operating characteristics of the sample Wasserstein-1 estimator $\hat{W}_1$ for the population value W_1 , reported in the original units of the effect modifier:

1. Bias: $\text{Mean}(\hat{W}_1) - W_1$
2. Root mean squared error (RMSE): $\sqrt{\text{Mean}[(\hat{W}_1 - W_1)^2]}$
3. Coverage probability: Proportion of 95% bootstrap CIs containing the true value
4. CI width: Mean width of bootstrap CIs, reflecting estimation precision

3.2 | Results

3.2.1 | Point Estimation and Coverage

Table 2 presents the bias of $\hat{W}_1$ across scenarios and sample sizes, in the original units of the effect modifier (scenarios are specified with $\sigma = 10$). The estimator showed positive bias under the null hypothesis (S1), with bias decreasing from 2.46 at $n = 50$ to 1.26 at $n = 200$. This positive bias arises because the non-negativity of W_1 yields $\hat{W}_1 > 0$ with probability 1 in finite samples, even when the true value equals zero ($F_1 = F_2$ almost everywhere). The estimator remains consistent ($\hat{W}_1 \rightarrow 0$ as $n \rightarrow \infty$ under the null), but the convergence is slow near the parameter space boundary.

TABLE 2 Bias of $\hat{W}_1$ by scenario and sample size, in the original units of the effect modifier.

Scenario	True W_1	$n = 50$	$n = 100$	$n = 200$
S1 (Null)	0.00	2.463	1.759	1.260
S2 (0.2σ)	2.00	0.990	0.472	0.173
S3 (0.5σ)	5.00	0.174	0.051	0.020
S4 (1.0σ)	10.00	0.034	−0.009	0.004
S5 (Scale)	3.99	0.829	0.389	0.189
S6 (Skew)	12.15	0.382	0.210	0.116
S7 (Loc+Scale)	5.85	0.624	0.326	0.181

Bias is defined as $\text{Mean}(\hat{W}_1) - W_1$, in the units of the effect modifier (scenarios specified with $\sigma = 10$). Results based on 10,000 replications with $B = 2,000$ bootstrap resamples.

For the non-null scenarios away from the boundary (S3–S7), bias was small relative to each scenario’s true W_1 and diminished with sample size; at $n = 200$ it was at most 0.189 (S5), a small fraction of the corresponding W_1 (Figure 2(A)). For S4 (1.0σ location shift, $W_1 = 10$), bias was negligible (0.034 at $n = 50$, 0.004 at $n = 200$), confirming accurate estimation even for large distributional differences. For scenarios near the boundary (S1, S2), positive bias was more pronounced and diminished slowly with sample size—a well-known phenomenon for non-negative statistics where sampling variability inflates estimates when the true value is close to zero.




FIGURE 2 Estimation properties of the sample Wasserstein-1 estimator $\hat{W}_1$, in the original units of the effect modifier, across scenarios (S1–S7) and sample sizes ($n = 50, 100, 200$). Panel (A): bias, with dashed horizontal line at zero. Panel (B): coverage probability of 95% percentile bootstrap CIs, with dashed line at the nominal 0.95 level; S1 coverage is structurally zero because the true value lies at the parameter space boundary (see Section 3.2). Panel (C): mean CI width in units of W_1 .

Table 3 presents coverage probabilities of the 95% bootstrap confidence intervals. For the non-null scenarios, coverage was near-nominal at $n \geq 100$: S3 (0.948–0.951), S4 (0.944–0.948), S6 (0.942–0.946), and S7 (0.936–0.945). S4 (1.0σ location shift) showed stable coverage across all sample sizes, with no degradation at larger n , confirming satisfactory bootstrap performance even for large distributional differences. S5 (scale) showed moderate undercoverage at $n = 50$ (0.863) that improved to near-nominal at $n = 200$ (0.945). S2 (small location shift) showed undercoverage at $n = 50$ (0.703) due to its near-boundary true value, improving to 0.947 at $n = 200$.

Coverage for S1 (null) is shown as zero in Figure 2(B) because its true ρ lies at the parameter space boundary. The non-negativity constraint produces persistent positive bias that prevents the finite-sample confidence interval from containing the true value of zero. This does not indicate estimator failure—the estimator is consistent—but rather that sample sizes well beyond 200 per region would be needed for reliable inference when the true distributional difference is negligible.

We also evaluated BCa (bias-corrected and accelerated) confidence intervals in preliminary experiments. BCa intervals showed consistently lower coverage than percentile intervals across all scenarios, particularly for near-boundary scenarios. The BCa correction overcorrects for $\hat{W}_1$ because the statistic is bounded below by zero, causing the acceleration parameter to distort the quantile adjustment. Based on these findings, we adopt percentile bootstrap as the primary inference method.

Monte Carlo standard errors for all reported coverage probabilities were below 0.005 at 10,000 replications, ensuring that observed differences across scenarios and sample sizes reflect genuine operating characteristic differences rather than simulation noise.

TABLE 3 Coverage probability of 95% percentile bootstrap confidence intervals.

Scenario	$n = 50$	$n = 100$	$n = 200$
S2 (0.2σ)	0.703	0.904	0.947
S3 (0.5σ)	0.950	0.951	0.948
S4 (1.0σ)	0.942	0.944	0.948
S5 (Scale)	0.863	0.917	0.945
S6 (Skew)	0.946	0.942	0.946
S7 (Loc+Scale)	0.928	0.936	0.945

Coverage under S1 (Null) is shown as zero in Figure 2(B) but omitted from this table for brevity; the structural zero arises from the non-negativity constraint at the parameter space boundary.

For S4, coverage was near-nominal across all sample sizes (0.942–0.948), confirming satisfactory performance even for large distributional differences. The undercoverage for S2 at small sample sizes (0.703 at $n = 50$) and S5 (0.863 at $n = 50$) reflects boundary proximity effects that diminish with increasing n .

3.2.2 | Estimation Precision

The mean CI width pattern across scenarios and sample sizes is shown in Figure 2(C). Table 4 presents the RMSE and mean CI width across scenarios, in the original units of the effect modifier. RMSE decreased consistently with increasing n , at the $O(n^{-1/2})$ rate expected from Wasserstein plug-in theory; for S3, for example, it fell from 1.83 at $n = 50$ to 0.98 at $n = 200$. Mean CI width narrowed proportionally, with CIs at $n = 100$ for the non-null scenarios spanning roughly 3.6–7.0 in units of W_1 . In the clinical calibration framework, this CI width propagates directly to $\Delta_{\max}$ via $\Delta_{\max} = L_{\text{clinical}} \cdot W_1$ (Section 2.4). For example, with $L_{\text{clinical}} = 0.01/\text{yr}$ (representative of age in the GUSTO-I application), a $\hat{W}_1$ CI width of 4.96 years (S3 at $n = 100$) corresponds to a $\Delta_{\max}$ CI width of $0.01 \times 4.96 = 0.050$ (5.0%pt)—a level of precision sufficient for meaningful clinical calibration.

TABLE 4 RMSE and mean CI width of $\hat{W}_1$ by scenario and sample size, in the original units of the effect modifier.

Scenario	RMSE			Mean CI Width		
	$n = 50$	$n = 100$	$n = 200$	$n = 50$	$n = 100$	$n = 200$
S1 (Null)	2.63	1.87	1.34	4.19	2.94	2.07
S2 (0.2σ)	1.60	1.13	0.86	4.67	3.57	2.84
S3 (0.5σ)	1.83	1.37	0.98	6.22	4.96	3.77
S4 (1.0σ)	1.98	1.42	1.00	7.52	5.45	3.89
S5 (Scale)	1.64	1.11	0.77	5.65	4.10	2.97
S6 (Skew)	2.61	1.86	1.29	9.74	7.04	5.03
S7 (Loc+Scale)	2.00	1.44	1.01	6.97	5.27	3.87

CI width is defined as mean of (upper – lower) across 10,000 replications, in the units of the effect modifier (scenarios specified with $\sigma = 10$).

Under the null hypothesis (S1), the positive bias produces CIs that are shifted upward from zero. At $n = 50$, the mean estimate was 2.46 with mean CI width 4.19, and at $n = 200$, the mean estimate decreased to 1.26 with mean CI width 2.07. This bias is relevant for clinical calibration. When the true distributional difference is zero, the estimator will suggest a small positive $\Delta_{\max}$. Practitioners should be aware of this conservative tendency, particularly at smaller sample sizes.

3.2.3 | Comparison with Standardized Mean Difference

Table 5 compares the sensitivity of $\hat{W}_1$ and SMD to different types of distributional differences. While both metrics respond to location shifts, SMD is insensitive to differences in variance and shape—the very types of distributional differences that may drive treatment effect heterogeneity through non-linear CATE functions.

TABLE 5 Sensitivity comparison: $\hat{W}_1$ versus SMD at $n = 100$.

Scenario	$\hat{W}_1$ (mean $\pm$ SD)	SMD (mean $\pm$ SD)	Implication
S3 (Location)	5.05 $\pm$ 1.36	0.50 $\pm$ 0.14	Both detect
S5 (Scale only)	4.38 $\pm$ 1.04	0.00 $\pm$ 0.14	Only $\hat{W}_1$ detects
S6 (Skew only)	12.36 $\pm$ 1.85	0.00 $\pm$ 0.14	Only $\hat{W}_1$ detects

SMD depends only on the difference in means, making it insensitive to differences in variance and skewness. $\hat{W}_1$ captures the full distributional difference. The comparison is one of detection: SMD is near zero for the pure scale (S5) and skew (S6) contrasts, whereas $\hat{W}_1$ is substantially positive, registering distributional differences that SMD cannot. $\hat{W}_1$ (in units of the effect modifier, here $\sigma = 10$) and SMD (dimensionless) are reported on different scales, so only the presence or absence of a signal, not the numerical magnitudes, should be compared across the two columns.

In summary, the simulation study across seven scenarios demonstrates that the $\hat{W}_1$ estimator provides reliable estimation with the following properties at sample sizes typical in MRCT regional subgroups:

1. Bias: Small relative to each scenario's W_1 at $n \geq 100$, with positive bias largest near the boundary (S1: 1.76 at $n = 100$) and negligible for the larger-shift scenarios (S3: 0.051, S4: -0.009 at $n = 100$).

2. Coverage: Near-nominal (0.92–0.95) at $n \geq 100$ for the non-null scenarios (S3–S7), including S6 (skew) and S7 (location + scale).
3. Precision: RMSE and CI width decrease at the $O(n^{-1/2})$ rate; at $n = 200$, RMSE ranged from 0.77 (S5) to 1.34 (S1) in units of W_1 .
4. Sensitivity: Captures location, scale, and skewness differences where SMD is limited to location only.

Scenario S3 (0.5σ location shift, true $W_1 = 5$) exemplifies the operating characteristics at clinically relevant effect sizes. Bias was negligible ($+0.051$ at $n = 100$), coverage was nominal (0.951), and the CI width (4.96 years) translated to a $\Delta_{\max}$ CI width of $0.01 \times 4.96 = 0.050$ (5.0%pt) when calibrated with $L_{\text{clinical}} = 0.01/\text{yr}$ (age-scale parameters representative of the GUSTO-I application)—a level of precision sufficient for informed regulatory deliberation.

These estimation properties ensure that $\hat{W}_1$, when combined with the clinical calibration framework of Section 2.4, provides $\Delta_{\max}$ estimates of sufficient quality for regulatory deliberation. We recommend $n \geq 100$ per region for reliable estimation and inference.

4 | APPLICATION: HYPOTHETICAL THROMBOLYTIC MRCT USING GUSTO-I DATA

We illustrate the per-EM Wasserstein-1 framework through a hypothetical scenario grounded in publicly available data from GUSTO-I, a large international trial in acute myocardial infarction (AMI). A sponsor is developing a novel thrombolytic agent (Drug T) for AMI and is planning a Phase 3 MRCT across multiple regions. At the planning stage, the sponsor uses IPD from GUSTO-I ($N = 40,830$; 16 regions)²¹ to assess distributional similarity of candidate effect modifiers across regions and to identify suitable pooling partners for a designated anchor region. GUSTO-I was not designed as an MRCT and its regions are anonymized.

4.1 | Scenario Design and Effect Modifier Selection

For this illustrative exercise, the sponsor identifies two candidate effect modifiers for Drug T's treatment effect in AMI: age and systolic blood pressure (SBP). Drug T is a novel agent. Its Phase 2 programme has suggested age and SBP as candidate effect modifiers on the basis of exploratory analyses and biological plausibility, but the precision of these analyses is insufficient for numerical estimation of L_{clinical} . At this planning stage the sponsor therefore has no quantitative prior on the clinical slope L_{clinical} for either candidate effect modifier, and class-level indirect evidence is of limited value as a numerical substitute. For age, the Fibrinolytic Therapy Trialists' (FTT) collaborative meta-analysis of existing fibrinolytic agents²⁰ concluded benefit is present irrespective of age and does not report a per-year CATE slope on the absolute mortality scale, so class transferability would not yield a defensible numerical prior for Drug T. For SBP, no class-level quantitative estimate of clinical slope is available at all. We therefore treat L_{clinical} as unknown a priori for both candidate effect modifiers, and apply the L^* reverse-calculation pathway (equation 9) to both.

The sponsor selects Region 8 as the anchor region—representing, for instance, a regulatory market where local efficacy evidence is required but sample size is limited. Region 8 comprises $n = 2,916$ patients in GUSTO-I, with baseline characteristics summarized in Table 6. The exercise of identifying Region 8's suitable pooling partners simulates the decision a sponsor would face when seeking regional efficacy assessment through pooling as described in ICH E17.³

TABLE 6 Region 8 baseline characteristics (GUSTO-I).

Candidate effect modifier	Mean	SD	Skewness	IQR _{pooled}
Age (years)	60.2	12.1	−0.17	17–18
SBP (mmHg)	132.4	22.9	0.20	30–34

$n = 2,916$. IQR_{pooled} ranges reflect variation across the 15 partner-region pairs.

4.2 | Distributional Assessment: Per-EM $\hat{W}_1$ Across 15 Partner Regions

We computed $\hat{W}_1$ between Region 8 and each of the remaining 15 regions for both candidate effect modifiers, expressed in the original measurement units (years for age, mmHg for SBP), with percentile bootstrap 95% confidence intervals ($B = 2,000$). Table 7 presents the results. Figure 3 displays these results as a forest plot with bootstrap CIs, with partners ordered by ascending $\hat{W}_{1,\text{age}}$ to aid visual comparison.

TABLE 7 Per-EM Wasserstein-1 distance $\hat{W}_1$ in original variable units (years for age, mmHg for SBP) with 95% percentile bootstrap CIs for Region 8 versus each partner region (GUSTO-I), listed by region number.

Partner	n	$\hat{W}_{1,\text{age}}$ (yr) [95% CI]	$\hat{W}_{1,\text{SBP}}$ (mmHg) [95% CI]
R1	2188	0.65 [0.51, 1.03]	3.61 [2.68, 4.93]
R2	2952	2.15 [1.57, 2.76]	0.95 [0.76, 1.82]
R3	2030	2.61 [1.96, 3.28]	3.35 [2.32, 4.65]
R4	2876	0.56 [0.37, 1.16]	2.60 [1.82, 3.75]
R5	1909	0.39 [0.35, 0.92]	3.37 [2.32, 4.59]
R6	1585	0.91 [0.63, 1.39]	3.37 [2.06, 4.77]
R7	3150	0.40 [0.30, 0.95]	5.07 [4.03, 6.22]
R9	3123	0.59 [0.40, 1.12]	6.80 [5.71, 8.01]
R10	1717	1.38 [0.79, 2.13]	4.00 [2.88, 5.39]
R11	2491	1.27 [0.76, 1.91]	6.44 [5.25, 7.71]
R12	4352	0.86 [0.47, 1.43]	6.40 [5.39, 7.46]
R13	2297	1.15 [0.76, 1.74]	2.40 [1.42, 3.72]
R14	3437	0.70 [0.42, 1.29]	4.36 [3.41, 5.44]
R15	2576	0.61 [0.40, 1.25]	4.72 [3.60, 5.94]
R16	1231	1.74 [1.16, 2.47]	4.28 [3.06, 5.71]

Point estimates with 95% percentile bootstrap CIs ($B = 2,000$, seed = 2026). Regions are listed by region number, and ordering carries no implicit ranking. Bootstrap CIs quantify estimation uncertainty. Narrow CIs lend greater confidence to individual $\hat{W}_1$ values, while overlapping CIs indicate partners whose relative position is not sharply resolved by the data.




FIGURE 3 Forest plots of $\hat{W}_1$ point estimates in original units (years for age, panel A; mmHg for SBP, panel B) with 95% percentile bootstrap CIs. Partners are ordered by ascending $\hat{W}_1$ within each panel; ordering and horizontal scale differ between panels.

The two candidate effect modifiers produce different patterns of distributional similarity. For age, partner-region $\hat{W}_{1,\text{age}}$ point estimates span from 0.39 years (R5) at the low end to 2.61 years (R3) at the high end, with most partners below 1.5 years.

For SBP, point estimates span from 0.95 mmHg (R2) to 6.80 mmHg (R9). Relative to their respective natural scales (age IQR ≈ 17 –18 years and SBP IQR ≈ 30 –34 mmHg in Region 8), the SBP differences are slightly larger in proportion than the age differences, and the ranking of partners by each candidate effect modifier differs accordingly.

Several partners fall in the middle of each ranking where uncertainty intervals overlap substantially. R16's $\hat{W}_{1,\text{age}}$ (1.74 years) is the 13th-smallest value with 95% CI [1.16, 2.47] years, placing it in a region where its relative position is not sharply determined by the data—the CI spans the range covered by several partners with smaller point estimates. R6's $\hat{W}_{1,\text{age}}$ (0.91 years) is the 9th-smallest value, while its $\hat{W}_{1,\text{SBP}}$ of 3.37 mmHg sits in the lower-middle of the SBP distribution with 95% CI [2.06, 4.77] mmHg, similarly spanning a broad region of the SBP ranking. These mid-rank cases illustrate the value of reporting bootstrap CIs alongside point estimates. The ranking itself is accompanied by an uncertainty interval that informs how confidently one partner can be placed above or below another.

An instructive contrast emerges between R2 and R9. R2 has the second-largest $\hat{W}_{1,\text{age}}$ (2.15 years) but the smallest $\hat{W}_{1,\text{SBP}}$ (0.95 mmHg), while R9 has the 4th-smallest $\hat{W}_{1,\text{age}}$ (0.59 years) but the largest $\hat{W}_{1,\text{SBP}}$ (6.80 mmHg). A ranking based on either candidate effect modifier alone would reach opposite conclusions about these two regions. This asymmetry underscores the necessity of evaluating all candidate effect modifiers jointly.

4.3 | Clinical Interpretation and Pooling Candidates

We require an upper bound on the clinical slope L_{clinical} for each candidate effect modifier; recall that L_{clinical} governs how strongly treatment effect varies with the modifier. For thrombolysis in acute myocardial infarction, class-level meta-analytic evidence (GUSTO-I,²¹ FTT collaborative meta-analysis,²⁰ and the GUSTO-I prognostic model²²) supports a conservative upper bound of approximately 10%pt per decade for age-related and 2%pt per 10 mmHg for SBP-related treatment effect modification on 30-day mortality, exceeding the empirical prognostic gradients in Lee et al.²² We adopt these as illustrative bounds: $L_{\text{age,UB}} = 1 \times 10^{-2}$ /yr and $L_{\text{SBP,UB}} = 2 \times 10^{-3}$ /mmHg.

For each partner region, we compare the region-specific L^* value (Table 8) with the corresponding plausible upper bound L_{UB} . A region is judged eligible for pooling on a given effect modifier if $L^* > L_{\text{UB}}$, meaning the clinical slope required for the observed distributional difference to produce $\Delta_{\text{clin}} = 1\%$ pt exceeds the plausible bound. On age, the eligible partners are R1, R4, R5, R6, R7, R9, R12, R14, R15 (nine regions with $L^*_{\text{age}} > 0.01$). On SBP, the eligible partners are R1, R2, R3, R4, R5, R6, R10, R13, R14, R15, R16 (eleven regions with $L^*_{\text{SBP}} > 0.002$); R7 ($L^*_{\text{SBP}} = 0.0020 = L_{\text{SBP,UB}}$) is excluded.

TABLE 8 Clinical interpretation (L^* at $\Delta_{\text{clin}} = 1\%$ pt) and joint eligibility for all 15 partner regions, listed by region number. L^* is the minimum clinical slope (per year for age, per mmHg for SBP) required for the observed distributional difference to produce a treatment effect difference of 1%pt on the absolute 30-day mortality scale. Eligibility is judged $L^* > L_{\text{UB}}$ with $L_{\text{age,UB}} = 1 \times 10^{-2}$ /yr and $L_{\text{SBP,UB}} = 2 \times 10^{-3}$ /mmHg.

Partner	$\hat{W}_{1,\text{age}}$ (yr)	$\hat{W}_{1,\text{SBP}}$ (mmHg)	L^* age (/yr)	L^* SBP (/mmHg)	Age elig.	SBP elig.	Joint
R1	0.65	3.61	0.0154	0.0028	✓	✓	✓
R2	2.15	0.95	0.0047	0.0105	—	✓	—
R3	2.61	3.35	0.0038	0.0030	—	✓	—
R4	0.56	2.60	0.0179	0.0038	✓	✓	✓
R5	0.39	3.37	0.0256	0.0030	✓	✓	✓
R6	0.91	3.37	0.0110	0.0030	✓	✓	✓
R7	0.40	5.07	0.0247	0.0020	✓	—	—
R9	0.59	6.80	0.0171	0.0015	✓	—	—
R10	1.38	4.00	0.0072	0.0025	—	✓	—
R11	1.27	6.44	0.0079	0.0016	—	—	—
R12	0.86	6.40	0.0116	0.0016	✓	—	—
R13	1.15	2.40	0.0087	0.0042	—	✓	—
R14	0.70	4.36	0.0143	0.0023	✓	✓	✓
R15	0.61	4.72	0.0164	0.0021	✓	✓	✓
R16	1.74	4.28	0.0057	0.0023	—	✓	—

Regions are listed by region number, and ordering carries no implicit ranking. ✓ indicates eligibility ($L^* > L_{\text{UB}}$); — indicates ineligibility.

A region is jointly eligible if it satisfies the criterion on both candidate effect modifiers, since pooling requires distributional compatibility across all evaluated modifiers. Six regions meet this AND condition: R1, R4, R5, R6, R14, R15. The remaining nine partners fail on at least one effect modifier—for example, R2 fails on age ($L_{\text{age}}^* = 0.0047$), R9 fails on SBP ($L_{\text{SBP}}^* = 0.0015$)—and are not recommended for pooling under these clinical assumptions.

Among the jointly eligible partners, R4 emerges as the leading single-pool candidate. R4 ranks 3rd-lowest on age ($\hat{W}_{1,\text{age}} = 0.56$ years) and 3rd-lowest on SBP ($\hat{W}_{1,\text{SBP}} = 2.60$ mmHg), balanced across both modifiers. R5 has the lowest age value ($\hat{W}_{1,\text{age}} = 0.39$ years) but a mid-rank SBP value ($\hat{W}_{1,\text{SBP}} = 3.37$ mmHg), an asymmetric profile. The framework supplies the quantitative inputs ($\hat{W}_1$ with bootstrap CIs in Table 7, partner-specific L^* in Table 8, and joint eligibility under the assumed bounds) to support sponsor judgment in consultation with clinical and regulatory advisors. Note that the GUSTO-I data were collected during 1990–1993; the application here serves as methodological illustration using publicly available IPD, not as an endorsement of these specific distributions for contemporary thrombolytic-trial planning.

5 | DISCUSSION

This paper addresses three methodological gaps identified in Section 2.1, demonstrating that the W_1 distance: (i) captures scale and skewness differences in regional effect modifier distributions that are invisible to standardized mean differences, as confirmed by simulation scenarios with location-preserving variance and shape contrasts (S5, S6); (ii) translates distributional distance into a clinically interpretable scale via the heterogeneity bound Δ_{max} (equation 8), providing a theoretical link to regional treatment effect heterogeneity that is absent both in empirical-distribution-function tests such as the Kolmogorov–Smirnov statistic and in representative-value distance methods such as Euclidean distance on region-level effect modifier summaries;⁷ and (iii) admits reliable percentile bootstrap inference for non-negligible distributional differences at moderate sample sizes ($n \geq 100$ per region, coverage 0.90–0.95 across S2–S7), while exhibiting characterized boundary behavior near the null that delimits the operational range of inference.

These methodological findings carry two interpretive implications for pooling decisions, illustrated by the GUSTO-I application as a worked example of the framework’s operation rather than a substantive epidemiological claim. First, when multiple candidate effect modifiers are under consideration, all candidates must be evaluated jointly. A partner ranked similar on one effect modifier may not be similar on another, as exemplified by the contrast between R2 and R9. Restricting evaluation to a subset of candidates risks selecting partners whose distributional differences on the omitted modifiers would later compromise consistency. Second, partner selection cannot rest on W_1 ranking alone. The required CATE sensitivity L^* that would translate a given W_1 into a clinically meaningful Δ_{max} varies markedly across partners and across candidate effect modifiers, so the same W_1 value carries different clinical weight depending on the partner-specific distributional differences. In the GUSTO-I example, six partners (R1, R4, R5, R6, R14, R15) emerged as jointly eligible because their L^* values exceeded the plausible bounds reviewed in Section 4; R4 stood out as the leading single-pool candidate due to its balanced ranking on both modifiers. Pooling judgments therefore require both distributional ranking and clinical calibration, applied across the full set of candidate effect modifiers.

These implications also clarify how the framework relates to the most concrete existing pooling recipe, that of Komiyama et al., who combine multiple effect modifiers into a single Lasso-weighted distance and cluster regions on the resulting distance matrix.⁷ The per-effect-modifier resolution adopted here matters precisely because partner rankings can reverse across modifiers: in the GUSTO-I example R2 and R9 each rank similar on one modifier but dissimilar on the other, so collapsing the modifiers into one aggregate distance can mask an incompatibility that the joint, per-modifier criterion of Section 4 retains. The two perspectives are nonetheless complementary, in two respects. First, the heterogeneity bound speaks directly to a question their account leaves open—whether a pooled-region treatment-effect estimate transfers to an individual member country. Within a pool whose members are mutually close on every candidate modifier, each pairwise regional difference in average treatment effect is bounded by $L_{\text{clinical}} \cdot W_1$ (equation 6), so the pooled estimate applies to each member country with a clinically interpretable and bounded error, rather than resting on the point estimate of the pool alone. Second, our distance-based calibration answers which partners are similar enough, a question distinct from how many pooled regions to form; subgroup-multiplicity considerations suggest the latter should remain small, typically no more than four,^{7,2} and the two constraints can be applied in sequence.

Section 2.1 identified three gaps in existing distributional measures. SMD captures only location. The KS statistic and related EDF tests lack a clinically interpretable scale and a theoretical link to treatment effect heterogeneity. KL divergence is asymmetric, can diverge under non-overlapping support, and requires density estimation impractical at small regional samples. The W_1 distance addresses all three. Where SMD captures only the mean, our simulation confirmed that scale (S5) and skewness (S6) differences produced near-zero SMD yet non-negligible W_1 . Where the KS statistic offers no clinical-scale translation, the

heterogeneity bound (equation 8) translates distributional distance into clinically meaningful units. Where KL divergence is asymmetric and prone to instability under limited sample overlap, W_1 is symmetric, finite for any pair of empirical distributions, and computable directly from CDFs without density estimation, while remaining amenable to bootstrap inference.

Two principal strengths follow from the dual-pathway clinical calibration. First, the framework adapts to where the trial sits on the spectrum of CATE sensitivity evidence, rather than imposing a fixed W_1 cutoff. At the MRCT planning stage, where L is typically unavailable for novel agents, the L^* reverse-calculation (equation 9) serves as the primary calibration tool. When L is available from prior evidence, the $\Delta_{\max}$ pathway (equation 8) provides a complementary calibration on the clinical scale. Second, the value of W_1 as a distributional measure, capturing scale and skewness differences invisible to SMD, holds regardless of whether L is available. Clinical calibration enhances but does not replace this distributional assessment.

For practice, we recommend computing W_1 with bootstrap CIs at moderate sample sizes (Section 3) for every candidate effect modifier across region pairs, translating each into $\Delta_{\max}$ (when L is estimable) or L^* at pre-specified Δ_{clin} values (when L is unknown), and reporting these alongside clinically relevant benchmarks (overall treatment effect, non-inferiority margin) to support deliberation rather than binary decisions. When multiple candidate effect modifiers are evaluated, sponsors may adopt a conservative approach based on the maximum $\Delta_{\max}$ (or smallest L^*) across all effect modifiers and region pairs, or a totality-of-evidence approach in which the full collection—together with bootstrap uncertainty and the reliability of each L estimate—informs holistic judgment. The choice depends on regulatory context.

For policy, the framework's compatibility with diverse data sources is a distinctive practical advantage. Because W_1 requires only baseline effect modifier distributions, regional data can be assembled from prior trials, registries, electronic health records, or other real-world evidence (RWE) sources—an increasingly relevant capability as regulatory agencies promote RWE use in drug development (ICH E6(R3)). The framework is particularly relevant for regulatory submissions requiring regional subpopulation consistency.^{4,5} Matsushima et al.'s⁴ PMDA workshop documented how regional imbalance in CRP-positive/MRI-negative status caused apparent inconsistency in the secukinumab MRCT, underscoring the practical consequences when effect modifier distributional differences are not assessed prospectively.

The W_1 distance is one input within a holistic regulatory evaluation; biological plausibility, clinical relevance, and benefit–risk balance extend beyond what any single distributional measure can quantify.^{3,4} The framework has further limitations. It treats continuous effect modifiers marginally, leaving categorical and multivariate extensions for future work. Being non-negative, W_1 exhibits positive bias and zero bootstrap coverage when the true value lies at or near zero (S1; Section 3.2),²⁴ so we recommend $n \geq 100$ per region and cautious interpretation of estimates near the null. Clinical calibration depends on the CATE sensitivity L , which may be unavailable for novel agents (handled by the L^* pathway) or transferred from prior evidence (a clinical assumption); the framework supplies quantitative inputs but prescribes no cutoffs. Finally, as with any covariate-based approach, unmeasured effect modifiers may contribute to heterogeneity beyond what the index captures.

Beyond the scope and methodological extensions noted above, two further directions warrant investigation. The upstream identification of which baseline characteristics constitute relevant effect modifiers is itself a non-trivial problem deserving dedicated treatment; regression-based relevance weighting such as the Lasso offers one front-end for this selection,⁷ after which the present framework quantifies distributional distance on the retained modifiers. Bias-correction methods tailored to the boundary behavior of W_1 at small samples could extend the operational range of inference. The principal contribution of this paper is to supply a standardized, clinically calibrated tool where none was previously available. The W_1 distance, combined with bootstrap confidence intervals and clinical calibration via $\Delta_{\max}$ or L^* , transforms the previously qualitative judgment of “similar enough” into a reproducible, quantitative basis for sponsor judgment, filling the methodological gap left by ICH E17 and enabling evidence-based, clinically grounded pooling decisions.

AUTHOR CONTRIBUTIONS

[Author 1] conceived the research question and led the project. [Author 2] developed the mathematical framework and conducted the simulation study. [Author 3] contributed to the application example and manuscript preparation. All authors reviewed and approved the final manuscript.

ACKNOWLEDGMENTS

The authors thank [colleagues] for helpful discussions on ICH E17 implementation.

FINANCIAL DISCLOSURE

None reported.

CONFLICT OF INTEREST

The authors declare no potential conflict of interests.

DATA AVAILABILITY STATEMENT

R code for computing the Wasserstein-1 distance and reproducing the simulation study is available at [repository URL]. GUSTO-I individual patient data are available from the original investigators subject to applicable data sharing policies; the dataset is described in the original publication.²¹

REFERENCES

- Chen J, Quan H, Binkowitz B, et al. Assessing consistent treatment effect in a multi-regional clinical trial: a systematic review. *Pharm Stat*. 2010;9(3):242–253. doi: 10.1002/pst.438
- Quan H, Li M, Chen J, et al. Assessment of consistency of treatment effects in multiregional clinical trials. *Drug Inf J*. 2010;44(5):617–632. doi: 10.1177/009286151004400515
- ICH E17 Expert Working Group. General Principles for Planning and Design of Multi-Regional Clinical Trials (E17). tech. rep., International Council for Harmonisation; Geneva, Switzerland: 2017.
- Matsushima N, Otsuba Y, Aoi Y, et al. Summary report of a public workshop: case studies of multi-regional clinical trial incorporating concept of the ICH E17 guideline. *Clin Pharmacol Ther*. 2024;115(5):965–970. doi: 10.1002/cpt.3163
- Song J, Ji C, Chen M, Luo X, Quan H, Chen J. Basic Considerations for Data Pooling Strategy in Multi-Regional Clinical Trials (MRCTs). *Ther Innov Regul Sci*. 2025;59(2):359–364. doi: 10.1007/s43441-025-00744-8
- Long M, Wu H, Liu X, Chen J. Basic Considerations for the Consistency Evaluation Based on ICH E17 Guideline. *Ther Innov Regul Sci*. 2025;59(2):328–336. doi: 10.1007/s43441-024-00737-z
- Komiyama O, Hiro S, Matsuoka N, Yamamoto H. Pooling Strategies. In: Li G, Binkowitz B, Wang W, Quan H, Chen J., eds. *Simultaneous Global New Drug Development: Multi-Regional Clinical Trials after ICH E17*, Chapman and Hall/CRC, 2024
- Austin PC. An introduction to propensity score methods for reducing the effects of confounding in observational studies. *Multivariate Behav Res*. 2011;46(3):399–424. doi: 10.1080/00273171.2011.568786
- Villani C. *Optimal Transport: Old and New*. Springer, 2009
- Vallender SS. Calculation of the Wasserstein distance between probability distributions on the line. *Theory Probab Appl*. 1974;18(4):784–786. doi: 10.1137/1118101
- Panaretos VM, Zemel Y. Statistical aspects of Wasserstein distances. *Annu Rev Stat Appl*. 2019;6:405–431. doi: 10.1146/annurev-statistics-030718-104938
- Gibbs AL, Su FE. On Choosing and Bounding Probability Metrics. *Int Stat Rev*. 2002;70(3):419–435. doi: 10.1111/j.1751-5823.2002.tb00178.x
- Barrio dE, Giné E, Matrán C. Central limit theorems for the Wasserstein distance between the empirical and the true distributions. *Ann Probab*. 1999;27(2):1009–1071. doi: 10.1214/aop/1022677394
- Efron B, Tibshirani RJ. *An Introduction to the Bootstrap*. New York: Chapman & Hall/CRC, 1993
- Sommerfeld M, Munk A. Inference for empirical Wasserstein distances on finite spaces. *J R Stat Soc Series B*. 2018;80(1):219–238. doi: 10.1111/rssb.12236
- Riley RD, Lambert PC, Abo-Zaid G. Meta-analysis of individual participant data: rationale, conduct, and reporting. *BMJ*. 2010;340:c221. doi: 10.1136/bmj.c221
- Fisher DJ, Carpenter JR, Morris TP, Freeman SC, Tierney JF. Meta-analytical methods to identify who benefits most from treatments: daft, deluded, or deft approach?. *BMJ*. 2017;356:j573. doi: 10.1136/bmj.j573
- VanderWeele TJ, Knol MJ. A Tutorial on Interaction. *Epidemiol Methods*. 2014;3(1):33–72. doi: 10.1515/em-2013-0005
- VanderWeele TJ. The interaction continuum. *Epidemiology*. 2019;30(5):648–658. doi: 10.1097/EDE.0000000000001054
- Fibrinolytic Therapy Trialists' (FTT) Collaborative Group. Indications for fibrinolytic therapy in suspected acute myocardial infarction: collaborative overview of early mortality and major morbidity results from all randomised trials of more than 1000 patients. *Lancet*. 1994;343(8893):311–322. doi: 10.1016/S0140-6736(94)91161-4
- The GUSTO Investigators. An international randomized trial comparing four thrombolytic strategies for acute myocardial infarction. *N Engl J Med*. 1993;329(10):673–682. doi: 10.1056/NEJM199309023291001
- Lee KL, Woodlief LH, Topol EJ, et al. Predictors of 30-day mortality in the era of reperfusion for acute myocardial infarction. Results from an international trial of 41,021 patients. *Circulation*. 1995;91(6):1659–1668. doi: 10.1161/01.cir.91.6.1659
- Wasserstein RL, Lazar NA. The ASA Statement on p-Values: Context, Process, and Purpose. *Am Stat*. 2016;70(2):129–133. doi: 10.1080/00031305.2016.1154108
- Andrews DWK. Inconsistency of the Bootstrap when a Parameter is on the Boundary of the Parameter Space. *Econometrica*. 2000;68(2):399–405. doi: 10.1111/1468-0262.00114

SUPPORTING INFORMATION

Additional supporting information may be found in the online version of the article at the publisher's website. Supporting materials include: complete R code for Wasserstein-1 distance computation, simulation scripts, and additional figures for realistic clinical scenarios.

How to cite this article: Author One, Author Two, and Author Three. Quantifying effect modifier similarity for regional pooling in multi-regional clinical trials. *Statist. Med*. 2025;00:1–18.

APPENDIX

A PROOFS AND MATHEMATICAL DETAILS

A.1 Non-negativity and Identity of Indiscernibles

Non-negativity is immediate: $W_1(F_1, F_2) = \int |F_1(x) - F_2(x)| dx \geq 0$, since the integrand is non-negative.

For the identity of indiscernibles, we show both directions. If $F_1 = F_2$, then $|F_1(x) - F_2(x)| = 0$ for all x , so $W_1(F_1, F_2) = 0$. Conversely, if $W_1(F_1, F_2) = \int |F_1(x) - F_2(x)| dx = 0$, then since $|F_1(x) - F_2(x)| \geq 0$ and is right-continuous, the integral can vanish only if $F_1(x) = F_2(x)$ almost everywhere, and by right-continuity of CDFs, $F_1 = F_2$ everywhere.

A.2 Asymptotic Properties

Under standard regularity conditions, the empirical Wasserstein-1 distance $\hat{W}_1 = W_1(\hat{F}_1, \hat{F}_2)$ is a consistent estimator of $W_1(F_1, F_2)$.¹³

In one dimension, the empirical W_1 distance converges at rate $O(n^{-1/2})$ without logarithmic correction.¹³ For two-sample estimation with $F_1 \neq F_2$ and finite second moments,

$$\sqrt{\frac{n_1 n_2}{n_1 + n_2}} (\hat{W}_1 - W_1(F_1, F_2)) \xrightarrow{d} N(0, \sigma_{W_1}^2), \quad (\text{A1})$$

where $n_1/(n_1 + n_2) \rightarrow \lambda \in (0, 1)$ and $\sigma_{W_1}^2$ depends on F_1 and F_2 through the limiting two-sample empirical process.¹³ This result requires $W_1 > 0$ (i.e., $F_1 \neq F_2$). At the boundary $F_1 = F_2$, the parameter lies on the edge of the parameter space and standard asymptotics do not apply, consistent with the degenerate bootstrap coverage observed under the null (S1; see Section 5).

The bootstrap provides valid inference for the Wasserstein distance under mild conditions.¹⁵ The percentile bootstrap confidence interval achieves asymptotically correct coverage for non-degenerate distributions.

B R CODE FOR W_1 COMPUTATION

Listing 1 R function for computing the Wasserstein-1 distance with bootstrap confidence intervals.

```
compute_W1 <- function(x1, x2) {
  pooled <- c(x1, x2)
  all_vals <- sort(unique(pooled))
  F1 <- ecdf(x1); F2 <- ecdf(x2)
  w1 <- sum(diff(all_vals) *
    abs(F1(all_vals[-length(all_vals)]) -
      F2(all_vals[-length(all_vals)])))
  w1
}

W1_bootstrap <- function(x1, x2,
  B = 2000, conf = 0.95) {
  obs <- compute_W1(x1, x2)
  boot_vals <- replicate(B, {
    b1 <- sample(x1, replace = TRUE)
    b2 <- sample(x2, replace = TRUE)
    compute_W1(b1, b2)
  })
  alpha <- 1 - conf
  ci <- quantile(boot_vals,
    c(alpha/2, 1 - alpha/2), na.rm = TRUE)
  list(estimate = obs,
    ci_lower = ci[1], ci_upper = ci[2])
}
```